

Distilling Knowledge from Graph Convolutional Networks

Yiding Yang¹, Jiayan Qiu², Mingli Song³, Dacheng Tao², Xinchao Wang^{1†}

¹Department of Computer Science, Stevens Institute of Technology, USA

²School of Computer Science, Faculty of Engineering, The University of Sydney, Australia

³College of Computer Science and Technology, Zhejiang University, China

{yyang99, xwang135}@stevens.edu, jqiu3225@uni.sydney.edu.au,

brooksong@zju.edu.cn, dacheng.tao@sydney.edu.au

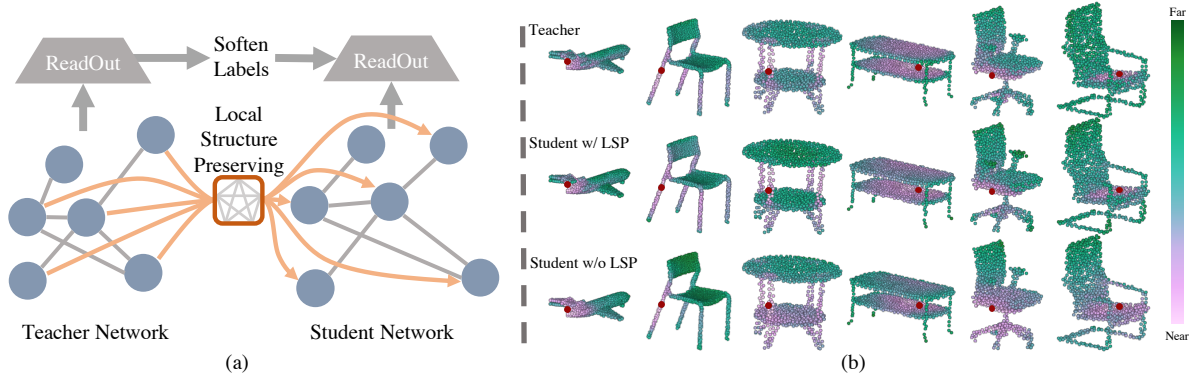


Figure 1: (a) Unlike existing knowledge distillation methods that focus on only the prediction or the middle activation, our method explicitly distills knowledge about how the teacher model embeds the topological structure and transfers it to the student model. (b) We display the structure of the feature space, visualized by the distance between the red point and the others on a **point cloud** dataset. Here, each object is represented as a set of 3D points. **Top Row**: structures obtained from the teacher; **Middle Row**: structures obtained from the student trained with the local structure preserving (LSP) module; **Bottom Row**: structures obtained from the student trained without LSP. Features in the middle and bottom row are obtained from the last layer of the model after training for ten epochs. As we can see, model trained with LSP learns a similar structure as that of the teacher, while the model without LSP fails to do so.

Abstract

Existing knowledge distillation methods focus on convolutional neural networks (CNNs), where the input samples like images lie in a grid domain, and have largely overlooked graph convolutional networks (GCN) that handle non-grid data. In this paper, we propose to our best knowledge the first dedicated approach to distilling knowledge from a pre-trained GCN model. To enable the knowledge transfer from the teacher GCN to the student, we propose a **local structure preserving module** that explicitly accounts for the topological semantics of the teacher. In this module, the local structure information from both the teacher and the student are extracted as distributions, and hence minimizing the distance between these distributions enables topology-aware knowl-

edge transfer from the teacher, yielding a compact yet high-performance student model. Moreover, the proposed approach is readily extendable to dynamic graph models, where the input graphs for the teacher and the student may differ. We evaluate the proposed method on two different datasets using GCN models of different architectures, and demonstrate that our method achieves the state-of-the-art knowledge distillation performance for GCN models. Code is publicly available at <https://github.com/ihollywhy/DistillGCN.PyTorch>.

1. Introduction

Deep neural networks (DNNs) have demonstrated their unprecedented results in almost all computer vision tasks. The state-of-the-art performances, however, come at the cost of the very high computation and memory loads, which in many cases preclude the deployment of DNNs on the

[†]Corresponding author.

edge side. To this end, knowledge distillation has been proposed, which is one of the main streams of model compression [14, 2, 38, 46]. By treating a pre-trained cumbersome network as the teacher model, knowledge distillation aims to learn a compact student model, which is expected to master the expertise of the teacher, via transferring knowledge from the teacher.

The effectiveness of knowledge distillation has been validated in many tasks, where the performance of the student closely approaches that of the teacher. Despite the encouraging progress, existing knowledge distillation schemes have been focusing on convolutional neural networks (CNNs), for which the input samples, such as images, lie in the grid domain. However, many real-life data, such as point clouds, take the form of non-grid structures like graphs and thus call for the graph convolutional networks (GCNs) [35, 10, 17, 12]. GCNs explicitly look into the topological structure of the data by exploring the local and global semantics of the graph. As a result, conventional knowledge distillation methods, which merely account for the output or the intermediate activation and omit the topological context of input data, are no longer capable to fully carry out the knowledge transfer.

In this paper, we introduce to our best knowledge the first dedicated knowledge distillation approach tailored for GCNs. Given a pre-trained teacher GCN, our goal is to train a student GCN model with fewer layers, or lower-dimension feature maps, or even a smaller graph with fewer edges. At the heart of our GCN distillation is the capability to encode the topological information concealed in the graph, which is absent in prior CNN-based methods. As depicted in Fig. 1, the proposed method considers the features of node as well as the topological connections among them, and hence provides the student model richer and more critical information about topological structure embedded by the teacher.

We illustrate the workflow of the proposed GCN knowledge distillation approach in Fig. 2. We design a local structure preserving (LSP) module to explicitly account for the graphical semantics. Given the embedded feature of node and the graph from teacher and student, the LSP module measures the topological difference between them and guides the student to obtain a similar topological embedding as the teacher does. Specifically, LSP first generates distribution for each local structure from both the student and the teacher, and then enforces the student to learn a similar local structure by minimizing the distance between the distributions. Furthermore, our approach can be readily extended to dynamic graph models, where the graphs are not static but constructed by the teacher and student model dynamically.

To see the distillation performance, we evaluate our method on two different tasks in different domains include Protein-Protein Interaction dataset for node classification and ModelNet40 for 3D object recognition. The models

used in these two tasks are in different architectures. Experiments show that our method consistently achieves the best knowledge distillation performance among all the compared methods, validating the effectiveness and generalization of the proposed approach.

Our contributions are summarized as following:

- We introduce a novel method for distilling knowledge from graph convolutional network. To the best of our knowledge, this is the first dedicated knowledge distillation method tailored for GCN models.
- We devise a local structure preserving (LSP) method to measure the similarity of the local topological structures embedded by teacher and student, enabling our method to be readily extendable to dynamic graph models.
- We evaluate the proposed method on two different tasks of different domains and on GCN models of different architectures, and show that our method outperforms all other methods consistently.

2. Related Work

Many methods have been proposed to distill the knowledge from a trained model and transfer it to a student model with smaller capacity [2, 38, 45]. Although there are many distillation strategies that do not only utilize the output [14] but also focus on the intermediate activation [31, 47, 16], they are all designed for the deep convolutional network with the grid data as input.

Our method, however, focuses on the graph convolutional networks, which handles a more general input that lies in the non-grid domain. To the best of our knowledge, this is the first attempt along this line. In what follows, we briefly review several tasks that are related to our method.

Knowledge Distillation. Knowledge Distillation (KD) is first proposed in [14], where the goal is distilling the knowledge from a teacher model that is typically large into a smaller model so that the student model can hold a similar performance as the teacher’s. In this method, the output of teacher is smoothed by setting a high temperature in the softmax function, which makes it contain the information of the relationship among classes. Besides the output, intermediate activation can also be utilized to better train a student network. FitNet [31] forces the student to learn a similar feature as teacher’s by adding an additional fully connected layer to transfer features of student model. [47] proposes a method to transfer the attention instead of the feature itself to get a better distillation performance. Moreover, NST [16] provides a method to learn a similar activation of the neurons. There are also many other methods in this field [23, 2, 38, 36, 24, 4, 34, 45, 5, 4], but none of them provide a solution that is suitable for GCN.

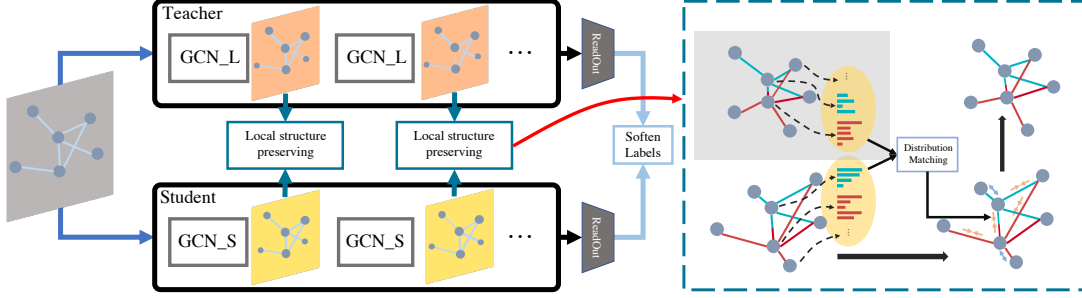


Figure 2: Framework of the proposed knowledge distillation method for GCNs. The local structure preserving module is the core of the proposed method. Given the feature maps and the graphs from the teacher and the student, we first compute the distribution of the local structure for each node and then match the distributions of the teacher with that of the student. The student model will be optimized by minimizing the difference of distribution among all the local structures.

Knowledge Amalgamation. Knowledge amalgamation [33, 22, 32, 44] aims to learn a student network from multiple teachers from different domains. The student model is trained as a multi-task model and learns from all the teachers. For example, a method of knowledge amalgamation is proposed in [43] to train a student model from heterogeneous-task teachers include scene parsing teacher, depth estimation teacher and surface-normal estimation teacher. The student model will have a backbone network trained from all the teachers’ knowledge and also several heads for different tasks trained from the corresponding teacher’s knowledge. MTZ [13] is a framework to compress multiple but correlated models into one model. The knowledge is distilled via a layer-wise neuron sharing mechanism. CFL [25] distills the knowledge by learning a common feature space, wherein the student model mimics the transformed features of the teachers to aggregate knowledge. Although many such methods are proposed, the models involved are usually limited within grid domain.

Graph Convolutional Network. In recent years, graph convolutional network [9, 6, 21, 27, 10, 42] has been proved to be a powerful model for non-grid data, which is typically represented as a set of nodes with features along with a graph that represents the relationship among nodes. The first GCN paper [17] shows that the GCN can be built by making the first-order approximation of spectral graph convolutions. A huge amount of methods have been proposed to make the GCN more powerful. GraphSAGE [12] gives a solution to make the GCN model scalable for huge graph by sampling the neighbors rather than using all of them. GAT [35] introduces the attention mechanism to GCN to make it possible to learn the weight for each neighbors automatically. [15] improves the efficient of training by adaptive sampling. In this paper, instead of designing a new GCN, we focus on how to transfer the knowledge effectively between different GCN models.

3D Object Recognition. One of the setups for 3D object recognition is predicting the label of the object given a

set of 3D points belong to it [28, 30]. Deep learning based methods [39, 11, 7, 20, 29, 26] outperform the previous methods that based on hand-crafted feature extractors [1, 3]. Moreover, GCN based methods [8, 19, 37, 18, 41], which can directly encode the structure information from the set of points, become one of the most popular directions along this line. Graph in these methods are typically obtained by connecting the k nearest points, where the distance is measured in the original space [18] or the learned feature space [37].

3. Method

In this section, we first give a brief description about the GCN followed by the motivation of the proposed knowledge distillation method that is based on the observation of the fundamental mechanism of GCN. We then provide the details about the local structure preserving (LSP) module, which is the core of our proposed method. Moreover, we explore the different choices of the distance functions used in the LSP module. Finally, we give the scheme to extend LSP to the dynamic graph models.

3.1. Graph Convolutional Network

Unlike the traditional convolutional networks that take grid data as input and output the high-level features, the input of graph convolutional networks can be non-grid, which is more general. Such non-grid input data is typically represented as a set of features $\mathbb{X} = \{x_1, x_2, \dots, x_n\} \in \mathbb{R}^F$, and a directed/undirected graph $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$. For example, in the task of 3D object recognition, we can set x_i as the 3D coordination and \mathcal{E} as the set of the nearest neighbors.

Given the input X and \mathcal{G} , the core operation of graph convolutional network is shown as:

$$x'_i = \mathcal{A}_{j:(j,i) \in \mathcal{E}} h_{\theta}(g_{\phi}(x_i), g_{\phi}(x_j)), \quad (1)$$

where h_{θ} is a function that considers the features in pair wise, g_{ϕ} is a function to map the features into a new space, \mathcal{A} is the strategy of how to aggregate features from the neighbors and get the new feature of the center node i .

There are many choices of function h , function g and the aggregation strategies. Take the graph attention network [35] as an example. it can be formulated as

$$x'_i = \sum_{j:j,i \in \mathcal{E}} \frac{e^{MLP_1(x_i || x_j)}}{\sum_{j:(j,i) \in \mathcal{E}} (e^{MLP_1(x_i || x_j)})} MLP_2(x_i), \quad (2)$$

where the function g is designed as a multilayer perceptron, function h is another multilayer perceptron that takes pair of nodes as input and predict the attention between them. The aggregation strategy is weighted summing all the features of neighbors according to the attention after normalization.

3.2. Motivation

The motivation of the proposed method is based on the fundamental of graph convolutional network. As shown in Eq. 1, the aggregation strategy (\mathcal{A}) plays an important role in embedding the features of nodes [21, 35], which is learned during the training process. We thus aim to provide the student the information about the function that the teacher has learned. However, it is challenging to distill knowledge that exactly represents the aggregation function and transfer it to the student. Instead of distilling the aggregation function directly, we distill the outcomes of such function: the embedded topological structure. The student can then be guided by matching the structure embedded by itself and that embedded by the teacher. We will show in the following sections how to describe the topological structure information and distill it to the student.

3.3. Local Structure Preserving

For the intermediate feature maps of a GCN, we can formulate it as a graph $\mathcal{G} = \{\mathcal{V}, \mathcal{E}\}$ and a set of features $\mathbb{Z} = \{z_1, z_2, \dots, z_n\} \in \mathbb{R}^F$, where n is the number of nodes, F is the dimension of the feature maps. The local structure can be summarized as a set of vectors $\mathbb{LS} = \{LS_1, LS_2, \dots, LS_n\}$, $LS_i \in \mathbb{R}^d$, where d is the degree of the center node i of the local structure. Each element of the vector is computed by

$$LS_{ij} = \frac{e^{SIM(z_i, z_j)}}{\sum_{j:(j,i) \in \mathcal{E}} (e^{SIM(z_i, z_j)})}, \quad (3)$$

$$SIM(z_i, z_j) = ||z_i - z_j||_2^2. \quad (4)$$

where SIM is a function that measures the similarity of the given pair of nodes, which can be defined as the euclidean distance between the two features. There are also many other advanced functions can be used here, which we will give more details in the following section. We take an exponential operation and normalize the values across all the nodes that point to center of the local structure. As a result, for each node i , we can obtain its corresponding local structure representation $LS_i \in \mathbb{R}^d$ by applying Eq. 3.

Notice that for different center node, their local structure representation may be in different dimension, which is depended on it's local graph.

In the setting of knowledge distillation, we are given a teacher network as well as a student one, where the teacher network is trained and fixed. We first provide here a local structure preserving strategy under the situation that both these two networks take the same graph as input but with different layers and dimension of embedded features. For the dynamic graph models, where the graph can be changed during the optimization process, we will give a solution in the section 3.5.

Given the intermediate feature maps, we can compute the local structure vectors for both the teacher and the student networks, which are denoted as \mathcal{LS}^s and \mathcal{LS}^t . For each center node i , the similarity of the local structure between the student's and the teacher's can be computed as

$$\mathcal{S}_i = D_{KL}(\mathcal{LS}_i^s || \mathcal{LS}_i^t) = \sum_{j:(j,i) \in \mathcal{E}} LS_{ij}^s \log\left(\frac{LS_{ij}^s}{LS_{ij}^t}\right), \quad (5)$$

where the Kullback Leibler divergence is adopted.

A smaller \mathcal{S}_i means a more similar distribution of the local structure. Thus, we compute the similarity of the distributions over all the nodes of the given graph and obtain the local structure preserving loss as

$$\mathcal{L}_{LSP} = \frac{1}{N} \sum_{i=1}^N \mathcal{S}_i. \quad (6)$$

The total loss is formulated as:

$$\mathcal{L} = \mathcal{H}(p_s, y) + \lambda \mathcal{L}_{LSP} \quad (7)$$

where y is the label and p_s is the prediction of the student model, λ is the hyperparameter to balance these two losses and \mathcal{H} represents cross entropy loss function that is also adopted by many other knowledge distillation methods [31, 16, 47].

3.4. Kernel Function

The similarity measurement function shown in Eq. 4 makes a strong assumption about the feature space that the similarity between pair of nodes is proportional to their euclidean distance, which is typically not the truth. We thus apply the kernel trick, which is very useful to map the vectors to a higher dimension and compute the similarity, to address this problem.

Kernel tricks are widely used in the traditional statistical machine learning methods. In that situation, the original feature vector will be mapped to a higher dimension by an implicit function φ . The similarity of the feature vector will then be computed as the inner product of the two mapped vectors as $\langle \varphi(z_i), \varphi(z_j) \rangle$. By adopting the kernel function,




Figure 3: Handling models with a dynamic graph, where the graph may be updated during the training process. We address this problem by first adding virtual edges according to the union of the two graphs from the student model and teacher model. The local structure preserving module can be then applied to the new graph directly.

we can compute the above two steps together without knowing the expression of φ .

There are several choices of the kernel functions. Three of the most common used kernel functions are linear function (Linear), polynomial kernel function (Poly) and radial basis function (RBF) kernel:

$$K(z_i, z_j) = \begin{cases} (z_i^T z_j + c)^d & \text{Poly} \\ e^{-\frac{1}{2\sigma^2} \|z_i - z_j\|^2} & \text{RBF} \\ z_i^T z_j & \text{Linear} \end{cases} \quad (8)$$

In this paper, we adopt and compare these three kernel functions and also the L_2 Norm. For the polynomial kernel function, d and c are set to two and zero respectively. For the RBF, σ is set to one.

3.5. Dynamic Graph

While the above method can be applied to the GCN models that with a fixed graph as input, it lacks flexibility for dynamic graph models. In the setup of dynamic graph model, both the feature of the nodes and the connection among them can be changed. Take the DGCNN model [37] as an example, the graph is initially constructed according to the 3D coordination of the input nodes/points and will be reconstructed once a new feature of nodes are obtained.

The advantage of dynamic graph is that the graph can represent the topological connection in a learned feature space rather than always in the initial feature space. However, such dynamic graph method can cause a problem to the local structure preserving module we mentioned above. The left of Fig. 3 is a case of the intermediate graphs generated by the DGCNN model. For each layer, the graph is

constructed by finding the K closet nodes for each node, where K is the hyperparameter. Directly computing the local structure vector using above method is meaningless because the distributions will come from a total different order of nodes.

We proposed a strategy to deal with such kind of situation by adding virtual edges to the graph in both the teacher model and student one. As shown in Fig. 3, given the two graphs constructed by the teacher and the student, which typically do not hold the same distribution of edges, we obtain the union of edges $\mathcal{E}_{\bullet i}^u = \{(j, i) : (j, i) \in \mathcal{E}^t | (j, i) \in \mathcal{E}^s\}$ for each center node i . A similar local structure vector can be obtained by replacing \mathcal{E} with \mathcal{E}^u in Eq. 3. Notice that although we adding the virtual edges for both the graphs, we only use them for local structure preserving module. The models still use their original graph to aggregate and update the feature of nodes. By considering the union structure of these two graphs, the generated local structure vectors now involve the nodes with the same distribution, which makes it possible for the student network to learn the topological relationship learned by the teacher.

Such strategy not only make it possible to compare the embedded local structure with different distribution of neighbors but also compare between local structure with different size of neighbors. It means that we can distill the knowledge from a teacher model with a large K to the student model with smaller K .

4. Experiments

In this section, we first give a brief description about the comparison methods. Then, we provide the experimental setup include the datasets we used, the GCN models we adopted and the details for each comparison methods. Notice that our goal is not to achieve the state-of-the-art performance in each dataset or task but to transfer as much as information from the teacher model the student one. This can be measured by the performance of the student model when the same teacher pre-trained model is involved.

We adopt two datasets in different domains, One is the protein-protein interaction (PPI) [48] dataset where the graphs come from the human tissues. This is a common used dataset for node classification [35, 12]. Another dataset is ModelNet40 [40] that contains point clouds come from the CAD models, which is a common used dataset for 3D object analysis [28, 30, 19, 37].

We also evaluate the proposed knowledge distillation method on GCN models with different architectures. Specifically, for the PPI dataset, GAT [35] model is adopted that takes fixed graph as input. For the ModelNet40 dataset, DGCNN [37] model with dynamic graph is adopted. We show in the experiments that our proposed method achieve the state-of-the-art distillation performance under various setups.

4.1. Comparison Methods

Since there is no knowledge distillation method designed for GCN models, we implement three knowledge distillation methods that can be used directly in a GCN model includes KD method [14], FitNet method [31] and attention transfer method (AT) [47]. Neuron selectivity transfer [16] is also one of the knowledge distillation methods for traditional convolutional networks. We leave it out since it makes an assumption that the size of the feature maps of student’s and teacher’s should be in the same size, which is not the case of our setup for the student and teacher models. Besides all the methods, we also set a baseline method which is training the student model with the original loss.

The summary of the comparison methods is as follow:

- **KD method [14]** is the first attempt for distilling knowledge from a teacher network. It utilizes a soften labels generated by the teacher network as an additional supervision. The intuition behind this method is that the soften label contains the similarity information among the classes learned by the teacher network. Since this method only relies on the output, it is suitable for most kinds of models.
- **FitNet method [31]** does not only make use of the output of the teacher model but also consider the intermediate feature maps. This method based on an assumption that the feature of the teacher model can be recovered from the feature of the student when the student is well trained. It introduces an additional mapping function to map the feature of student’s to that of the teacher’s and compute the L_2 distance between the mapped feature and the teacher’s feature.
- **Attention transfer method (AT) [47]** provides another way to transfer the knowledge in attention domain. In this method, the student model is forced to focus on the similar spatial areas like the teacher does, which is achieved by adding an L_2 loss between their attention maps. The attention map can be obtained from the feature maps and keep in the same size without consideration of the different channels.

4.2. Node Classification

In the node classification task, we are given the input nodes with associated features and also the graph. The goal is to generate the embedded feature for each node such that nodes with different classes can be separated. We adopt protein-protein interaction (PPI) dataset that contains 24 graphs corresponding to different human tissues. We follow the same dataset splitting protocol, wherein 20 graphs are used for training, two graphs are used for validation and another two graphs are used for testing. The average number of nodes for each graph in this dataset is 2372 and each

node has an average degree of 14. The dimension of input feature of the nodes is 50 and the number of class is 121.

For this dataset and task, we adopt GAT model for both the teacher model and student model. Since this is a multi-label task, where each node can belong to more than one classes, the binary cross entropy loss is adopted. The architecture of these two models are shown as follow:

Model	Layers	Attention heads	Hidden features
Teacher	3	4,4,6	256,256,121
Student	5	2,2,2,2,2	68,68,68,68,121

Table 1: Summary of the teacher and student models used on the PPI dataset for node classification. The student network is deeper than the teacher but with lower dimension of hidden features.

Model	Params	RunTime	Training	F1 Score
Teacher	3.64M	48.5ms	1.7s/3.4G	97.6
Student_Full	0.16M	41.3ms	1.3s/1.2G	95.7
Student_KD [14]	-	-	-	-
Student_AT [47]	0.16M	41.3ms	1.9s/1.4G	95.4
Student_FitNet [31]	0.16M	41.3ms	2.4s/1.6G	95.6
Student_LSP (Ours)	0.16M	41.3ms	2.0s/1.5G	96.1

Table 2: Node classification results on the PPI dataset. Teacher model used on this dataset is a GAT model with three hidden layers. full means the student model is trained with the ground truth labels without the teacher model.

The results are shown in Tab. 2. Params represents the total number of parameters; RunTime is the inference time for one sample and Training is the training time/GPU memory usage for one iteration, which is measured in a Nvidia 1080Ti GPU. The optimizer, learning rate, weight decay and training epochs are set to Adam, 0.005, 0 and 500 respectively for all the methods. All other hyperparameters for each method are tuned to obtain the best results on the validation set. Specifically, for AT method, the attention is computed by as $\sum_{i=1}^C |F_i|$; For ours, the kernel function is set to RBF and λ is set to 100. Notice that the loss function does not involve the softmax function. Therefore, the KD method, which is used by setting a high temperature in a softmax function, is not suitable here.

As what can be seen, all the knowledge distillation methods expect ours fail to provide a positive influence to the student model and lead to a drop of the performance respect to the performance of model trained with the original loss. Our method, thanks to the ability to transfer the local structure information to the student model, provides a positive influence and lead to a student model with the best performance among all the comparison methods.

4.3. 3D Object Recognition

We adopt ModelNet40 [40] dataset for the 3D object recognition task, wherein the object is represented by a set of points with only the 3D coordination as features. There

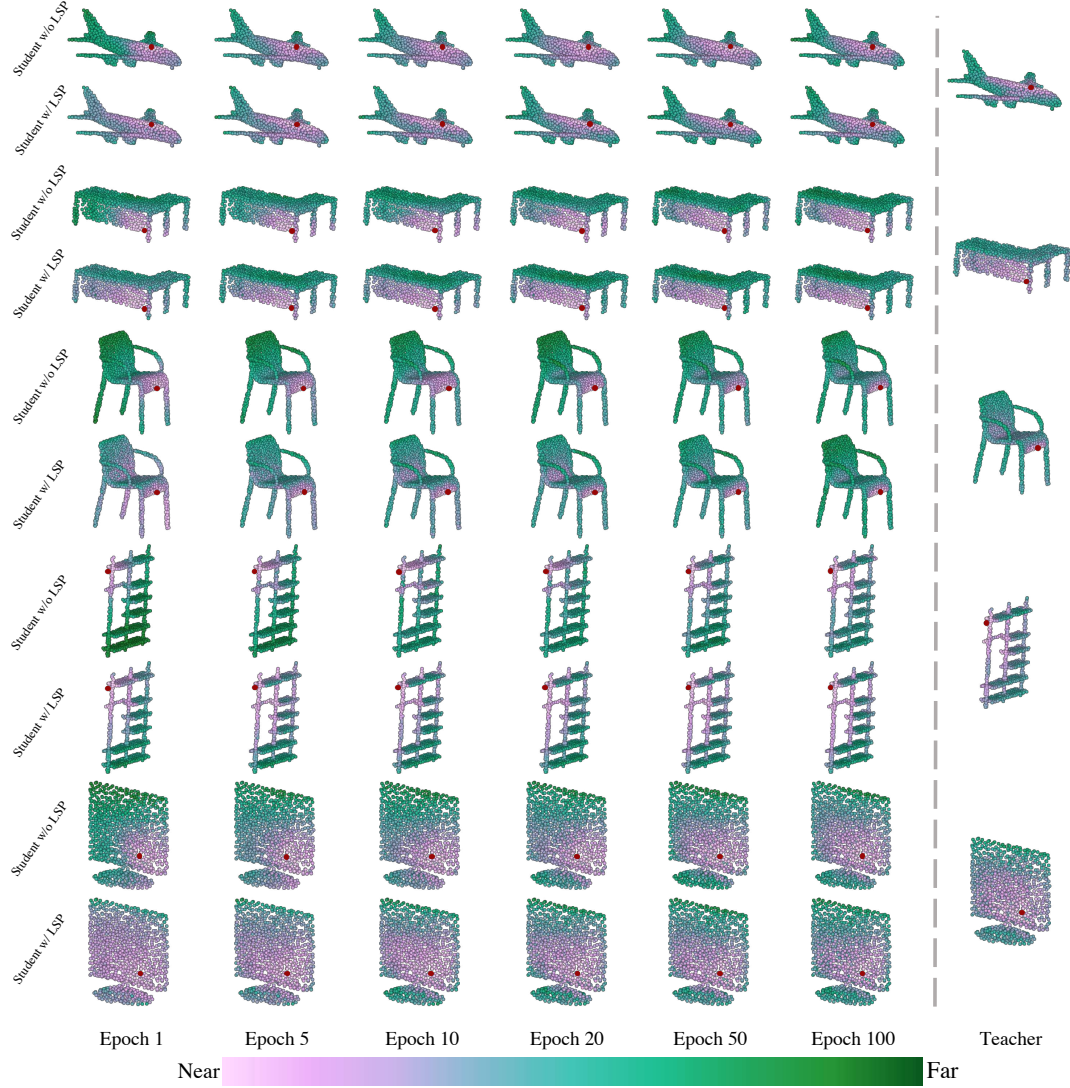


Figure 4: Visualization of the structure of the learned feature space, depicted by the distance between the red point and others. The features are extracted from the last layer of models. From left to right, we show the structure from the model trained in epoch 1, 5, 10, 20, 50, and 100. The **rightmost column** are obtained from the teacher model. For each object, the **lower row** is obtained from student model trained with our proposed method, and the **upper row** is obtained from student model trained with cross entropy loss. Our proposed knowledge distillation method guides the student model to embed the local structure as the teacher model does, leading to a similar structure in the very early training stage.

are 40 classes on this dataset and objects in each class come from the CAD models.

In this experiment, the architecture of both the teacher and student models is the same as DGCNN [37]. DGCNN is a dynamic graph convolutional model that makes both the advantages of the PointNet [28] and graph convolutional network. The graph here is constructed according to the distance of the points in their feature space. Since the feature space is different in different layers and different training stages, the graph is also changed, making it a dynamic graph

convolutional model.

The teacher model is under the same setup as the original paper: it has five graph convolutional layers followed by two fully connected layers. The student model has four graph convolutional layers with fewer feature map channels followed by one fully connected layers. The size of graph used in the student model, which is determined by the number of neighbors (K) when constructing the graph, is also smaller than the teacher's. The details of these two models are summarized in Tab. 3.

Model	Layers	Feature map's size	MLPs	K
Teacher	5	64,64,128,256,1024	512,256	20
Student	4	32,32,64,128	256	10

Table 3: Summary of the teacher and student models used on the ModelNet40 dataset. The student network is with less layers, fewer channels, and smaller input graphs.

The results are shown in Tab. 4. For the KD method, α is set to 0.1, the same as the original paper. The optimizer, learning rate, momentum and training epoch are set to SGD, 0.1, 0.9, 250 respectively for all the comparison methods. The other hyperparameters for each method are turned to obtain the best accuracy on the validation set. For our method, the kernel function is set to RBF and λ is set to 100.

We can see from the results that both KD method and AT method can boost the performance of the student model. Our method, thanks to the ability to learn the structure information from dynamic graph, generates the best student model for both the accuracy and mean class accuracy. The student model with only fix percent parameters can achieve a similar performance as the teacher model, which shows the generalization ability of the proposed method for different tasks and different architectures.

Model	Params	RunTime	Training	Acc	mAcc
Teacher	1.81M	8.72ms	30s/4.2G	92.4	89.3
Student_Full	0.1M	3.31ms	12s/1.4G	91.2	87.5
Student_KD [14]	0.1M	3.31ms	15s/1.4G	91.6	88.1
Student_AT [47]	0.1M	3.31ms	21s/1.7G	91.6	87.9
Student_FitNet [31]	0.1M	3.31ms	28s/2.4G	91.1	87.9
Student_LSP (Ours)	0.1M	3.31ms	29s/2.2G	91.9	88.6

Table 4: 3D object recognition results on ModelNet40. Teacher network used here is a DGCNN model with four graph convolutional layers.

4.4. Structure Visualization

In order to provide an intuitive understanding of the proposed method, we visualize the structure of the learned feature space during the process of optimization, which is represented by the distance among points in the object. As shown in Fig. 4, the student model trained with the proposed method (shown in the lower row for each object) can learn a similar structure as the teacher model very quickly in the early training stage. It can partially explain why the proposed method can generate a better student model.

4.5. Ablation and Performance Studies

To thoroughly evaluate our method, we provide here ablation and performance studies include the influence of different kernel functions as well as the performance for different student model configurations. All the experiments for this section are conducted on the ModelNet40 dataset.

Different Kernel Functions. We test all the three different kernel functions as Eq. 8 and also the naive L_2 Norm. The results are shown in Tab. 5. All the functions can provide the positive information to get a better student model and RBF works the best.

Model	Acc	mAcc
LSP w/ L_2 Norm	91.4	88.3
LSP w/ Polynomial function	91.7	87.7
LSP w/ RBF	91.9	88.6
LSP w/ Linear function	91.5	87.8

Table 5: Performance with different kernel functions. RBF achieves the overall best performance.

Different Model Configurations. We provide here the experiments to evaluate the trade-off between the model complexity and the performance when training with our proposed method. Specifically, we change the student model by adding more channels to each layer, adding one more graph convolutional layers and adding one more fully connected layers. Notice that student model with more channels holds almost the same accuracy performance as the teacher model but with less many parameters.

Model	Params (M)	RunTime (ms)	Acc	mAcc
W/ more Channels	0.44	3.83	92.3	88.7
W/ more Layers	0.14	4.26	92.1	89.2
W/ more MLPs	0.30	3.67	91.8	88.6

Table 6: Performance with different model configurations to evaluate the trade-off between performance and model-size/run-time.

5. Conclusion

In this paper, we propose a dedicated approach to distilling knowledge from GCNs, which is to our best knowledge the first attempt along this line. This is achieved by preserving the local structure of the teacher network during the training process. We represent the local structure of the intermediate feature maps as distributions over the similarities between the center node of the local structure and its neighbors, so that preserving the local structures are equivalent to matching the distributions. Moreover, the proposed approach can be readily extended to dynamic graph models. Experiments on two datasets in different domains and on two GCN models of different architectures demonstrate that the proposed method yields state-of-the-art distillation performance, outperforming existing knowledge distillation methods.

Acknowledgement

This work is supported by Australian Research Council Projects FL-170100117, DP-180103424 and Xinchao Wang’s startup funding of Stevens Institute of Technology.

References

- [1] Mathieu Aubry, Ulrich Schlickewei, and Daniel Cremers. The wave kernel signature: A quantum mechanical approach to shape analysis. In *2011 IEEE international conference on computer vision workshops (ICCV workshops)*, pages 1626–1633. IEEE, 2011.
- [2] Jimmy Ba and Rich Caruana. Do deep nets really need to be deep? In *Advances in neural information processing systems*, pages 2654–2662, 2014.
- [3] Michael M Bronstein and Iasonas Kokkinos. Scale-invariant heat kernel signatures for non-rigid shape recognition. In *2010 IEEE Computer Society Conference on Computer Vision and Pattern Recognition*, pages 1704–1711. IEEE, 2010.
- [4] Hanting Chen, Yunhe Wang, Chang Xu, Zhaohui Yang, Chuanjian Liu, Boxin Shi, Chunjing Xu, Chao Xu, and Qi Tian. Data-free learning of student networks. *arXiv preprint arXiv:1904.01186*, 2019.
- [5] Tianqi Chen, Ian Goodfellow, and Jonathon Shlens. Net2net: Accelerating learning via knowledge transfer. *arXiv preprint arXiv:1511.05641*, 2015.
- [6] Jian Du, Shanghang Zhang, Guanhang Wu, José MF Moura, and Soumya Kar. Topology adaptive graph convolutional networks. *arXiv preprint arXiv:1710.10370*, 2017.
- [7] Yi Fang, Jin Xie, Guoxian Dai, Meng Wang, Fan Zhu, Tiantian Xu, and Edward Wong. 3d deep shape descriptor. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2319–2328, 2015.
- [8] Yifan Feng, Haoxuan You, Zizhao Zhang, Rongrong Ji, and Yue Gao. Hypergraph neural networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 3558–3565, 2019.
- [9] Matthias Fey, Jan Eric Lenssen, Frank Weichert, and Heinrich Müller. Splinecnn: Fast geometric deep learning with continuous b-spline kernels. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 869–877, 2018.
- [10] Justin Gilmer, Samuel S Schoenholz, Patrick F Riley, Oriol Vinyals, and George E Dahl. Neural message passing for quantum chemistry. In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pages 1263–1272. JMLR. org, 2017.
- [11] Kan Guo, Dongqing Zou, and Xiaowu Chen. 3d mesh labeling via deep convolutional neural networks. *ACM Transactions on Graphics (TOG)*, 35(1):3, 2015.
- [12] Will Hamilton, Zhitao Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems*, pages 1024–1034, 2017.
- [13] Xiaoxi He, Zimu Zhou, and Lothar Thiele. Multi-task zipping via layer-wise neuron sharing. In *Advances in Neural Information Processing Systems*, pages 6016–6026, 2018.
- [14] Geoffrey Hinton, Oriol Vinyals, and Jeffrey Dean. Distilling the knowledge in a neural network. In *NIPS Deep Learning and Representation Learning Workshop*, 2015.
- [15] Wenbing Huang, Tong Zhang, Yu Rong, and Junzhou Huang. Adaptive sampling towards fast graph representation learning. In *Advances in Neural Information Processing Systems*, pages 4558–4567, 2018.
- [16] Zehao Huang and Naiyan Wang. Like what you like: Knowledge distill via neuron selectivity transfer. *arXiv preprint arXiv:1707.01219*, 2017.
- [17] Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907*, 2016.
- [18] Loic Landrieu and Martin Simonovsky. Large-scale point cloud semantic segmentation with superpoint graphs. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 4558–4567, 2018.
- [19] Yangyan Li, Rui Bu, Mingchao Sun, Wei Wu, Xinhan Di, and Baoquan Chen. Pointcnn: Convolution on x-transformed points. In *Advances in Neural Information Processing Systems*, pages 820–830, 2018.
- [20] Yangyan Li, Sören Pirk, Hao Su, Charles R Qi, and Leonidas J Guibas. Fpnn: Field probing neural networks for 3d data. In *Advances in Neural Information Processing Systems*, pages 307–315, 2016.
- [21] Yujia Li, Daniel Tarlow, Marc Brockschmidt, and Richard Zemel. Gated graph sequence neural networks. *arXiv preprint arXiv:1511.05493*, 2015.
- [22] Iou-Jen Liu, Jian Peng, and Alexander Schwing. Knowledge flow: Improve upon your teachers. In *International Conference on Learning Representations*, 2019.
- [23] Yufan Liu, Jiajiong Cao, Bing Li, Chunfeng Yuan, Weiming Hu, Yangxi Li, and Yunqiang Duan. Knowledge distillation via instance relationship graph. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 7096–7104, 2019.
- [24] Yufan Liu, Jiajiong Cao, Bing Li, Chunfeng Yuan, Weiming Hu, Yangxi Li, and Yunqiang Duan. Knowledge distillation via instance relationship graph. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 7096–7104, 2019.
- [25] Sihui Luo, Xinchao Wang, Gongfan Fang, Yao Hu, Dapeng Tao, and Mingli Song. Knowledge amalgamation from heterogeneous networks by common feature learning. *arXiv preprint arXiv:1906.10546*, 2019.
- [26] Daniel Maturana and Sebastian Scherer. Voxnet: A 3d convolutional neural network for real-time object recognition. In *2015 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 922–928. IEEE, 2015.
- [27] Federico Monti, Davide Boscaini, Jonathan Masci, Emanuele Rodola, Jan Svoboda, and Michael M Bronstein. Geometric deep learning on graphs and manifolds using mixture model cnns. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 5115–5124, 2017.
- [28] Charles R Qi, Hao Su, Kaichun Mo, and Leonidas J Guibas. Pointnet: Deep learning on point sets for 3d classification and segmentation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 652–660, 2017.
- [29] Charles R Qi, Hao Su, Matthias Nießner, Angela Dai, Mengyuan Yan, and Leonidas J Guibas. Volumetric and

- multi-view cnns for object classification on 3d data. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 5648–5656, 2016.
- [30] Charles Ruizhongtai Qi, Li Yi, Hao Su, and Leonidas J Guibas. Pointnet++: Deep hierarchical feature learning on point sets in a metric space. In *Advances in neural information processing systems*, pages 5099–5108, 2017.
 - [31] Adriana Romero, Nicolas Ballas, Samira Ebrahimi Kahou, Antoine Chassang, Carlo Gatta, and Yoshua Bengio. Fitnets: Hints for thin deep nets. *arXiv preprint arXiv:1412.6550*, 2014.
 - [32] Chengchao Shen, Xinchao Wang, Jie Song, Li Sun, and Mingli Song. Amalgamating knowledge towards comprehensive classification. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 3068–3075, 2019.
 - [33] Chengchao Shen, Mengqi Xue, Xinchao Wang, Jie Song, Li Sun, and Mingli Song. Customizing student networks from heterogeneous teachers via adaptive knowledge amalgamation. *arXiv preprint arXiv:1908.07121*, 2019.
 - [34] Antti Tarvainen and Harri Valpola. Mean teachers are better role models: Weight-averaged consistency targets improve semi-supervised deep learning results. In *Advances in neural information processing systems*, pages 1195–1204, 2017.
 - [35] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Lio, and Yoshua Bengio. Graph attention networks. *arXiv preprint arXiv:1710.10903*, 2017.
 - [36] Haoyu Wang, Defu Lian, and Yong Ge. Binarized collaborative filtering with distilling graph convolutional networks. *arXiv preprint arXiv:1906.01829*, 2019.
 - [37] Yue Wang, Yongbin Sun, Ziwei Liu, Sanjay E Sarma, Michael M Bronstein, and Justin M Solomon. Dynamic graph cnn for learning on point clouds. *arXiv preprint arXiv:1801.07829*, 2018.
 - [38] Zhenyang Wang, Zhidong Deng, and Shiyao Wang. Accelerating convolutional neural networks with dominant convolutional kernel and knowledge pre-regression. In *European Conference on Computer Vision*, pages 533–548. Springer, 2016.
 - [39] Zhirong Wu, Shuran Song, Aditya Khosla, Fisher Yu, Linguang Zhang, Xiaoou Tang, and Jianxiong Xiao. 3d shapenets: A deep representation for volumetric shapes. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1912–1920, 2015.
 - [40] Zhirong Wu, Shuran Song, Aditya Khosla, Fisher Yu, Linguang Zhang, Xiaoou Tang, and Jianxiong Xiao. 3d shapenets: A deep representation for volumetric shapes. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1912–1920, 2015.
 - [41] Yaoqing Yang, Chen Feng, Yiru Shen, and Dong Tian. Foldingnet: Point cloud auto-encoder via deep grid deformation. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 206–215, 2018.
 - [42] Yiding Yang, Xinchao Wang, Mingli Song, Junsong Yuan, and Dacheng Tao. Spagan: Shortest path graph attention network. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI-19*, pages 4099–4105. International Joint Conferences on Artificial Intelligence Organization, 7 2019.
 - [43] Jingwen Ye, Yixin Ji, Xinchao Wang, Kairi Ou, Dapeng Tao, and Mingli Song. Student becoming the master: Knowledge amalgamation for joint scene parsing, depth estimation, and more. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 2829–2838, 2019.
 - [44] Jingwen Ye, Xinchao Wang, Yixin Ji, Kairi Ou, and Mingli Song. Amalgamating filtered knowledge: Learning task-customized student from multi-task teachers. *arXiv preprint arXiv:1905.11569*, 2019.
 - [45] Junho Yim, Donggyu Joo, Jihoon Bae, and Junmo Kim. A gift from knowledge distillation: Fast optimization, network minimization and transfer learning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 4133–4141, 2017.
 - [46] Xiyu Yu, Tongliang Liu, Xinchao Wang, and Dacheng Tao. On compressing deep models by low rank and sparse decomposition. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, July 2017.
 - [47] Sergey Zagoruyko and Nikos Komodakis. Paying more attention to attention: Improving the performance of convolutional neural networks via attention transfer. *arXiv preprint arXiv:1612.03928*, 2016.
 - [48] Marinka Zitnik and Jure Leskovec. Predicting multicellular function through multi-layer tissue networks. *Bioinformatics*, 33(14):i190–i198, 2017.